


What the bare-utterance labelling can bear



n=1,080 utterances, 270 per turn. Later messages are elliptical, and a single-utterance classifier has less to work with. Bars are 95% intervals. Confidence does not decline, so the drift is not the model degrading; part of it is length.